

In creative production, we constantly need to run quick-turnaround simulations, model training, cache-baking, neural texture generation, inference runs, asset processing, or heavy workloads that exceed local CPU and GPU compute power. **Spark Fuse** gives you exactly that: instant access to serious compute, with zero infrastructure to manage.



Spark Fuse: Fast and scalable compute for every job

Spark Fuse is Spark Cloud Studio's ephemeral GPU and CPU compute layer for creative technologists, developers, pipeline teams and ML engineers.

Submit a single Python function, Git repo, Docker image, executable, basic script, full project tree or even just a simple one-line command. Spark Fuse runs your job on demand, streams the logs, writes outputs directly back into your ShareSync file storage as they're created, and shuts down when your job is complete. You're only charged for the compute time your job used.



No manual provision



No idle machines



No manual file transfers

Two compute engines. Pick the one that fits the job.

InstantCompute

- ✔ Launches in 1-2 seconds
- ✔ GPU from \$0.49/hour
- ✔ CPU from \$0.19/hour
- ✔ Best for rapid iteration, debugging, inference tests, and anything where you need instant results

SmartCompute


- ✔ Launches in a few minutes
- ✔ GPU from \$0.29/hour
- ✔ CPU from \$0.09/hour
- ✔ Best for bursting cost-sensitive workloads, batch processing, checkpointed jobs, and interrupt-tolerant tasks

InstantCompute when you can't wait. SmartCompute when you can.
Same hardware. Same API. Same outputs.



Save up to 70% with SmartCompute for scaling workloads

How it works



Why choose Spark Fuse over other serverless GPU platforms?

Spark Fuse	Runpod/Modal
Unified ecosystem: One login for workstations, storage, and compute	Compute layer only
Flexible pricing: From \$0.09/hr for smart savings or \$0.19/hr for speed	Not optimized for specific tasks
File handling: Output directly back into your storage space in real-time.	Manual S3/bucket management

Additional benefits

- + **TPN-aligned security:** Built for secure production workflows, with all data encrypted end-to-end and fully silo'd to your organisation.
- + **Private by design:** Customer workloads run on isolated compute, with files governed by workspace permissions. Your data is your data.
- + **Built by industry professionals:** Based on production experience across workflows, storage, render, and pipeline infrastructure.

Guide pricing *

SmartCompute

NVIDIA T4 from \$0.29/hr
RTX PRO 6000 from \$1.19/hr

InstantCompute

NVIDIA T4 from \$0.49/hr
RTX PRO 6000 from \$2.89/hr

Pay only while your job runs.
No idle machines.
No subscription required.


Technical snapshot

- + RTX PRO 6000 Blackwell, L40S, A100, H100, and more
- + TPN compliant security architecture
- + Global compute availability
- + Multi-GPU configurations available by account and region
- + High-core CPU workloads supported
- + No limits on concurrent jobs
- + Cached model weights for repeat jobs
- + Live stdout and stderr logs
- + Secure Spark OAuth authentication
- + Outputs written directly to ShareSync

Spark connects everything

Together, along with **Spark Cloud Studio** offerings including workstations, secure file storage and an infinite render farm, **Spark Fuse** gives you a complete production-ready and scalable compute ecosystem, always ready to go.

- One login.
- One billing system.
- One platform.



Get your Spark Fuse key today

Spark Fuse is now in early access. Contact us now to run your first job.



Email us at ignite@sparkcloud.studio for your key and get started in minutes

<https://sparkcloud.studio>

* Rates vary by GPU type, region, current demand and availability. Detailed pricing estimates for all jobs are available through the Spark API.