

Sometimes you need to train a LoRA. Run a ComfyUI workflow as a batch. Fine-tune a model. Train a Gaussian splat. Test simulation settings. Bake neural textures. Roto a shot with SAM2. Upscale a sequence overnight. Run an inference job on a few thousand frames that far exceed local CPU and GPU compute power. **Spark Fuse** gives you exactly that: instant access to serious compute in the cloud, with zero infrastructure to manage. Fully pay-as-you-go. No subscriptions required.

Fast and scalable compute for every job

Run sims, LoRAs, inference jobs, batch scripts, asset prep, and pipeline tools without tying up local machines.

Two compute engines. Pick the one that fits the job.

InstantCompute	SmartCompute
✔ Launches in 1-2 seconds ✔ GPU from \$0.49/hour ✔ CPU from \$0.19/hour ✔ Best for rapid iteration, debugging, inference tests, and instant results	✔ Launches in a few minutes ✔ GPU from \$0.29/hour ✔ CPU from \$0.09/hour ✔ Best for bursting, batch processing, and interrupt-tolerant tasks



TPN-aligned security



Private by design



Built by industry professionals

A few ways Spark Fuse can transform your workflows

Department	Example Job	Processes	Output
AI/Lookdev	LoRA training, ComfyUI workflows, style tests	Training scripts, inference containers, model pipelines	LoRA weights, sample grids, reference sheets
Nuke/Comp	Batch comps, denoise, ML mattes, cleanup	Batch comps, denoise, ML mattes, cleanup	EXR, matte, , processed plates
Maya/Assets	Exports, conversions, validations, previews	Batch scripts, USD and alembic workflows, mayapy automations	USD, Alembic, reports, previews, FBX
Houdini/FX	Pyro, FLIP, Vellum, RBD wedges	hython, hbatch, Cache scripts, Dockerised sim tools	VDB, Alembic, USD, FBX, MOV
Pipeline/TD	Texture, metadata extraction, colour processing, QC	OpenImageIO, FFmpeg, OCIO, custom tooling	Processed assets, QC reports, proxies
Editorial	Proxy generation, transcodes, slates, burn-ins	RVIO tools, custom scripts, annotation plugins	ProRes, MP4, EXR, EDL, playlists, rough cuts

How it works



Submit your jobs through our simple Spark API. Spark Fuse then packages your job, selects and runs it on the right cloud GPU or CPU while giving you live logs so you can track progress in real time. Your job's results are written straight back into your private ShareSync storage.

Simple and easily accessible compute for the heavy tasks that usually clog workstations, overnight boxes, or ad-hoc pipeline machines.



Local training isn't sustainable. Use **Spark Fuse** to train and compare multiple LoRA variants faster while keeping approved outputs organised in ShareSync.

Guide pricing *

SmartCompute

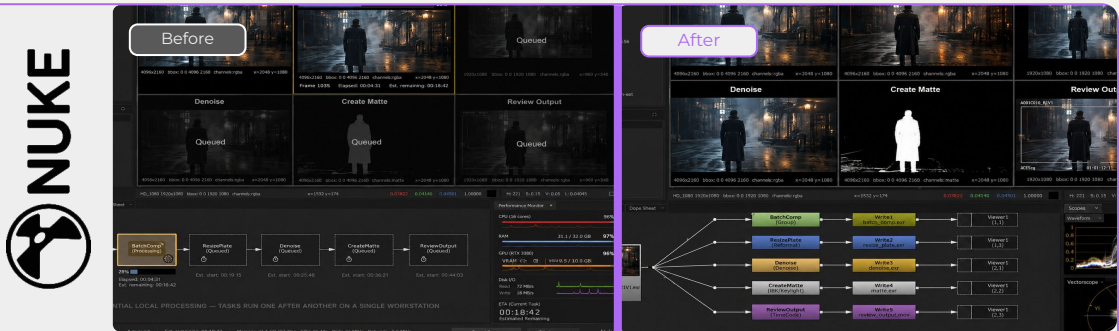
NVIDIA T4 from \$0.29/hr
RTX PRO 6000 from \$1.19/hr

InstantCompute

NVIDIA T4 from \$0.49/hr
RTX PRO 6000 from \$2.89/hr

Pay only while your job runs.

No idle machines.
No subscription required.



Offload comps, resizes, denoise scripts, matte generations, and automated workflows to free up local resources. Let **Spark Fuse** do all the grunt work for you.

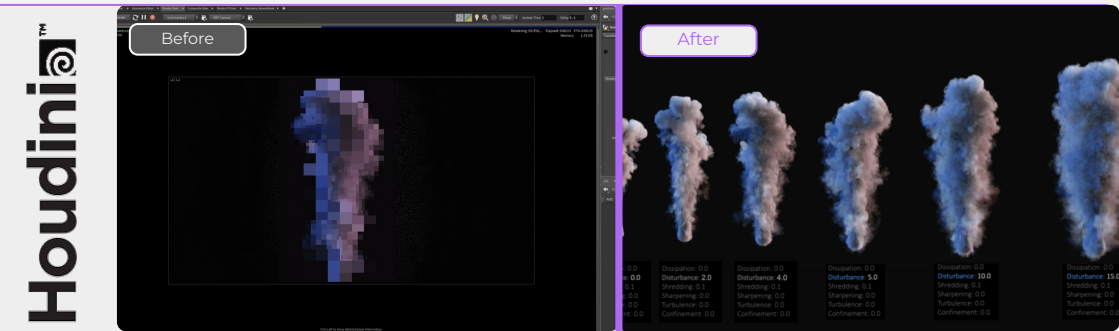
Spark connects everything

Together, along with **Spark Cloud Studio** offerings including workstations, secure file storage and an infinite render farm, **Spark Fuse** gives you a complete production-ready and scalable compute ecosystem, always ready to go.

One login.

One billing system.

One platform.



Iterate, automate and set off multiple sets of wedged simulations with **Spark Fuse**, so you can work freely in your interactive session.

Get your Spark Fuse key today

Spark Fuse is now in early access.
Contact us now to run your first job.



Email us at ignite@sparkcloud.studio
for your key and get started in minutes

<https://sparkcloud.studio>